

HIMANSHU SUTHAR

AI Engineer Intern | Multi-Agent Systems · RAG · LLM Integration · Prompt Engineering · Full-Stack (Python & React)
sutharindustry@gmail.com • +91 9660484162 • LinkedIn • GitHub • Portfolio

PROFILE

B.Tech CSE (AI & ML) student shipping production-grade AI systems. Built multi-agent fraud detection (CrewAI, 4 agents, tool-calling), LLM-powered email automation (n8n + FastAPI), a RAG pipeline with vector search, and a clinical AI platform with 98.3% accuracy. Hands-on with OpenAI, Anthropic, and Grok APIs. Comfortable turning ambiguous problems into working features and shipping independently.

CORE SKILLS

Languages: Python, JavaScript, TypeScript, SQL
AI/LLM: CrewAI, OpenAI & Anthropic APIs, Grok API (xAI), n8n, LangChain, Prompt Engineering, Multi-Agent Orchestration, Tool-Calling, Response Evaluation
RAG & Retrieval: Vector Search, Embeddings, FastAPI RAG Pipelines, Knowledge Base Retrieval, Chunking & Similarity Search
Web & Backend: React, JavaScript (ES6+), Node.js, FastAPI, REST APIs, Postgres, Supabase, Responsive UI
ML & Data: Scikit-learn, PyTorch, TensorFlow, Pandas, Matplotlib, Seaborn, AUC/ROC Evaluation, Cohort & Funnel Analysis
Engineering & Cloud: Git & GitHub, Docker, Google Cloud (Cloud Run, BigQuery), Agile/SDLC, JIRA, Unit Testing

EXPERIENCE

- AI Engineer Intern**

Jul 2026 – Present

NOSO Labs (YC S25) · San Francisco Bay Area

 - Building AI agents and workflows for technician and sales use cases in the trades — prototyping LLM-powered features for field teams, dispatch, and business operators.
 - Developing RAG systems using company and customer knowledge bases; integrating LLMs with CRM, call, and field service systems.
 - Improving prompt quality and model evaluation through structured feedback loops; creating internal tools for data processing and customer onboarding.
- AI Engineer**

Apr 2026 – May 2026

The AI Signal

 - Built a local-first RAG backend (FastAPI + vector search) optimized for low-latency inference — document chunking, embedding generation, and similarity search pipeline.
 - Debugged performance bottlenecks through structured evaluation of chunk sizes and embedding models; authored architecture documentation for engineering handoffs.
- Data Analyst**

May 2026 – Present

3Skill

 - Analyzed 10,000+ records with Python and SQL — data validation, null-handling, deduplication — ensuring analytics correctness before stakeholder reporting.
 - Built and maintained stakeholder-facing dashboards, iterating on feedback in a fast-paced collaborative environment.
- AI Summer Intern**

Jul 2026

Mirai School of Technology

 - Built a Multi-Modal Visual Novel Engine integrating 3 AI APIs (Gemini, Pollinations, gTTS) with fault-tolerant error handling and dynamic JSON-based UI rendering.
 - Applied SDLC and Agile practices — requirements analysis, coding, unit testing — under senior developer mentorship.
- Data Science & Analytics Intern**

Jul 2026

Future Interns · Remote

 - Completed 3 analytics projects (Sales, Churn, Funnel) across 7,000+ records — KPI analysis, cohort analysis, churn prediction, and funnel optimization with 20+ visualizations.

KEY PROJECTS

Fraud Detection — CrewAI Multi-Agent System

GitHub

Python · CrewAI · OpenAI API · Tool-Calling · Statistical Anomaly Detection

- Built 4 specialized agents (Transaction Analyst, Anomaly Detector, Risk Assessor, Report Generator) working sequentially with shared context across the pipeline.
- Implemented custom tool-calling: Z-score anomaly detection, velocity fraud checks, geographic flagging, high-risk merchant detection.

Automated Email Classifier — n8n + LLM Workflow

GitHub

n8n · FastAPI · OpenAI API · Gmail/Slack/Calendar APIs · Prompt Engineering

- 14-node n8n workflow: polls inbox → LLM classifies into 6 categories with confidence scores → routes to Gmail labels, Slack alerts, calendar events, and email forwarding via Switch routing.
- Designed FastAPI endpoint with structured prompt engineering for JSON output — entity extraction, priority scoring, and classification accuracy iteration.

RAG Pipeline — Retrieval-Augmented Generation Backend

GitHub

Python · FastAPI · Vector Search · Embeddings · LLM APIs

- Built RAG pipeline — document chunking, embedding generation, similarity search, LLM-grounded response synthesis — optimized for low-latency inference.

SportChat — AI Sports Conversation App

GitHub

JavaScript · React · Grok API (xAI) · Prompt Engineering · REST APIs

- Full-stack chat app with real-time LLM-driven conversations — prompt engineering to ground responses in sports context for domain-specific accuracy.

Healthcare AI System — Clinical Intelligence Console

si14.oneapp.dev

Python · FastAPI · Scikit-learn · PyTorch · 8 ML Algorithms · REST APIs

- Trained 8 algorithms on clinical datasets — 98.3% accuracy, 0.997 AUC. Unified behind a single FastAPI backend serving 8 REST endpoints.

Cardiac Risk Atlas — Heart-Disease Risk Estimator

cra12.oneapp.dev

Python · Logistic Regression · Statistical Modeling · Data Imputation

- Logistic regression on Cleveland (n=1,025) and Framingham (n=4,240) datasets — AUC 0.93 with median imputation for missing values.

EDUCATION

B.Tech — Computer Science Engineering (AI & ML)

Aug 2025 – May 2029

Lovely Professional University (LPU)

Coursework: Data Structures & Algorithms · Software Engineering · Statistics · Machine Learning

CERTIFICATIONS & ACHIEVEMENTS

- 5-Day AI Agents Intensive Course — Google
- Software Engineering Job Simulation — Electronic Arts
- Data Visualisation: Empowering Business with Effective Insights — Tata Group
- NeuroHack 2026 Hackathon — Built CivicAI (Google Cloud, BigQuery, Cloud Run)